

# Link Prefetching: A Defense Against Website Fingerprinting on Tor \*

Vaibhav Sharma  
sharm361@umn.edu

Taejoon Byun  
taejoon@umn.edu

Elaheh Ghassabani  
ghass013@umn.edu

Se Eun Oh  
seoh@umn.edu

Department of Computer Science and Engineering  
University of Minnesota  
Minneapolis, MN 55454

## ABSTRACT

*This paper is written to get an A in the CSCI5271 course. PERIOD.*

*This content will be edited later.* We plan to explore the area of website fingerprinting in anonymization networks starting with the paper Website Fingerprinting in Onion Routing Based Anonymization Networks.

## Keywords

Website fingerprinting, anonymity, encrypted traffic, Tor

## 1. INTRODUCTION

Anonymizing networks are privacy technologies that provide a mechanism to anonymize internet communications so as to protect users from network eavesdroppers. Although such systems are able to hide the communication (including both routing information and content), an attacker is still able to obtain different information by analyzing the network traffic. Network analysis can provide very rich information about message length, timing, and frequency by which an attacker can easily identify the communicating parties, and therefore bypass an anonymizing system. This problem is known as Website Fingerprinting (WF) attack, where the adversary attempts to recognize the encrypted traffic patterns of specific web pages without using any other information [6, 7].

This report is organized as follows;

## 2. WEBSITE FINGERPRINTING ON TOR

This section explains the background about Tor, an anonymity network, and the website fingerprinting attack which can

neutralize the anonymity that Tor provides.

### 2.1 Tor

The Onion Router (Tor) is one of the most popular anonymizing networks, which distributes user's communications over several places on the Internet. The distribution helps user's final destination not to be linked to a single point. In other words, packets take a random path through several relays, instead of taking a direct route. These layers cover the user path such that it will not be possible to tell where the data comes from, or where it is going. Such a pathway in Tor is called a private network pathway. The Tor client application is responsible of creating this private pathway by incrementally building a circuit of encrypted connections. The circuit is constructed hop by hop, which means each individual relay is only aware of the previous and next hop. That is to say, individual relays never can identify the complete path [2, 3].

If an attacker can see both ends of the private pathway, Tor will fail. For example, suppose the attacker watches the Tor relay to which the user enters the network, and also watches the website he visits; In this case, the attacker can correlate volume and timing information on the two sides. To cope with this problem, Tor came up with the idea of "entry guards", where each Tor client randomly selects a few relays as entry points for her first hop. If those relays are not observed by an attacker, the user is secure. But, needless to say, there is always a probability of losing anonymity [2].

### 2.2 Website Fingerprinting

Tor does not provide protection against all anonymity problems. In addition WF attack is still able to bypass its privacy mechanism. In the case of Tor, WF attack would take place between the user and the Guard node, or at the Guard node itself. In general, attack models can be divided into two categories: "closed world" and "open world" scenarios. This section describes these attack models.

In the closed world model, the classifier can successfully recognize only web pages that it has already been trained on. Therefore, in this case, we deal with a small set of censored web pages. The open world scenario is when the ability of the adversary can recognize censored pages that might have not been seen before [1].

Figure 1 illustrates the attack model of website fingerprinting.

\*This report is submitted as a partial fulfillment of CSCI5271: Introduction to Security course.



Figure 1: Website Fingerprinting Attack Model

### 3. RELATED WORK

There are a decent number of works on bypassing anonymizing systems, especially Tor, on which this section provides an overview.

#### 3.1 Fingerprinting Attacks

Penchenko et al. [9] were among the first to report website fingerprinting attacks with reasonable accuracy on Tor. This paper provides a sufficient understanding of the feature set and classification framework required for this attack. [9] is a comprehensive study of WF on Tor, which considers both closed world and open world attack models.

#### 3.2 Defense Mechanisms

#### 3.3 Criticism

### 4. LINK PREFETCHING AS A DEFENSE

This section explains the concept of link prefetching and discusses a possible usage of it as a defense mechanism against fingerprinting attacks on Tor.

#### 4.1 Link Prefetching

Link prefetching is a HTML syntax that gives the web browser hints about which page the user is most likely to visit in a near future [4, 5]. The pages and resource to pre-fetch are specified in the web page so that the web browser can silently load them (or pre-fetch them) after an idle time. Since the pre-fetch happens only after the page is fully loaded, it does not sacrifice the loading time of the requested web page. Moreover, it can save the loading time for the pre-fetched pages and thus improve the user experience by caching the *future* contents. It was first suggested by Mozilla Foundation in 2003 and supported by most modern browsers nowadays.

The resources to pre-fetch can be simply specified in *HTML* using a *link* tag [8]. For example, a *link* tag `<link rel="prefetch" href="/page2.html">` tells the browser to pre-fetch a *html* file named *page2.html*. Resources other than a *HTML* web page can also be pre-fetched similarly using the same syntax. There are also some variations for different types of prefetching – DNS prefetching, which is specified as `<link rel="dns-prefetch" .>`, is supported by *Mozilla Firefox* and *Google Chrome*, and `<link rel="prerender" .>` also does the same job as *prefetch* in *Google Chrome* and *Microsoft Internet Explorer*.

Figure 2 illustrates how prefetching actually works in a browser (*Google Chrome*). This page is set up arbitrarily by the authors to demonstrate link prefetching, and contains a link pre-fetch tag that specifies to load a big image (the image on the left side labeled as *prefetch*). When the prefetching is off, this image shall be requested only when a user puts his mouse cursor on it (mouse over). However, it can be seen on the network timeline (right bottom) that

Table 1: Experiment design to answer *RQ1*

victim \ attacker	prefetch on	prefetch off
prefetch on	(1)	(2)
prefetch off	(3)	(4)

the image is pre-fetched right after loading the page. This is indicated by a long blue bar on the second row for the file named *Very-high...*, and it is long because the size of the file is relatively big (*3.5 MB*) that it took a relatively long time to download. Please also note that the time it took for loading the image when the user actually requested (by putting his mouse on it) was very short, because the image had already been pre-fetched that the browser merely loaded it from the cache (as shown in the *size* column on the fourth row).

#### 4.2 The Effect of Prefetching on Fingerprinting

The idea of studying the effect of prefetching on Tor network is not new [?].

Figure 3 shows an aspect of fingerprint we gathered for the first page of [www.wired.com](http://www.wired.com). It depicts the number of packets on the *y-axis* for each time interval (*500 ms*), where the *x-axis* represents the time in seconds. The solid line represents the number of packets throughout time when the pre-fetch feature was enabled, and the dashed line represents when the pre-fetch was disabled in the browser settings.

### 5. EXPERIMENTAL EVALUATION

Experimental result comes here.

#### 5.1 Research Question

1. **RQ1:** Does prefetching itself provide an extra degree of defense?
2. **RQ2:** Can prefetching be used as a browser-side defense mechanism?
1. We speculate that prefetching itself might provide extra defense because of the extra packets. It can also be the case however, prefetching websites are more vulnerable to fingerprinting because of the extra prefetch packets that shows distinct prefix.
2. This case is unlikely since prefetching is on by default. We assume that victims will more likely be using Tor under the default setting.
3. This case is what we are most curious about, whether a victim can confuse an attacker by simply turning prefetching setting off of his browser.
4. This case simulates a situation where a victim is loading any other websites that does not prefetch any resource. This can be used as a comparison case.

#### 5.2 Effects of Pre-Fetching on Fingerprinting

In this section, we will write about our experiments. We are planning to conduct two sets of experiments. If we consider the network traffic, the number of packages go upstream depends on the number of pre-fetching requests, and

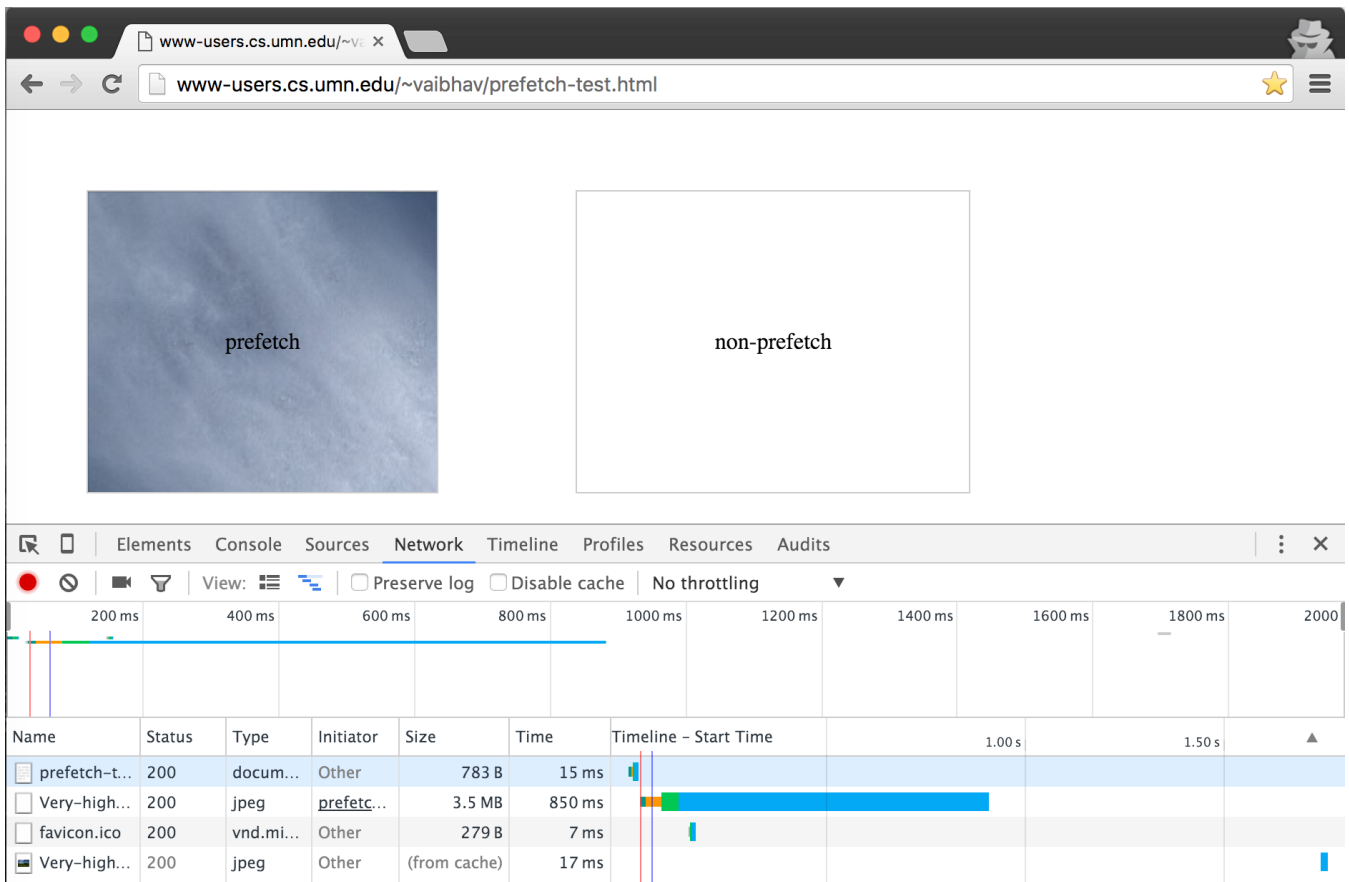


Figure 2: Network timeline showing pre-fetch




Figure 3: hi

the number of downstream packages coming depends on the size of resources that should be pre-fetched. Therefore, it is obvious that pre-fetching would affect the fingerprint of the traffic of a particular website. *should be completed.*

### 5.3 Investigate Pre-Fetching Effects on top 60 Popular Websites

We are running experiments to see how pre-fetching affects the websites' fingerprints. After doing some search on top popular websites, we put together a small crawler by

which we learnt that only around 60 of all 6000 websites are use pre-fetching mechanism. We are capturing traffic of these websites in two different modes: 1) with enabled pre-fetching, and 2) with disabled pre-fetching. We are working on feature extraction, and about to decide which classifiers to use for the learning phase. Ultimately, we plan to conduct two sets of experiments. One sort of experiment is to compare two series of the captured packets and find the accuracy number with the help of a classifier, by which our goal is to provide an evidence to see if pre-fetching really affects fingerprints of websites. So, if the result will be positive, we will perform another set of experiment, which kind of simulates a sub set of those 60 websites. Then, we will see how (altering) the size of pre-fetching affects fingerprinting attacks/ defense mechanisms.

### 5.4 Effect of Pre-Fetching Packets Size on Fingerprinting Attacks

Here, we will explain our second experiment. We will simulate a sub set of webpages we investigated in the previous experiment. Then, we will equip them with a mechanism so that they can affect the downstream traffic and finally their fingerprint. Then, we will analyze the result to see how this idea contributes to the effectiveness of attacks and defense techniques.

## 6. DISCUSSION

We'll find what to discuss later on after finishing the experiment section.

## 7. CONCLUSIONS

What we have done in this work is so cool and awesome. What you may criticize will all be put here as "future work".

This section will conclude the result of our experiments. Finally we will provide some evidence to show how prefetching affect fingerprinting attacks. Based on our result, we are planing to suggest some defense mechanisms.

## 8. ACKNOWLEDGMENTS

The authors appreciate Professor Stephen McCamant for telling geeky jokes in classes all the time. This is a research project for CSCI5271, University of Minnesota.

## 9. REFERENCES

- [1] A critique of website traffic fingerprinting attacks. <https://blog.torproject.org/blog/critique-website-traffic-fingerprinting-attacks>, 2015. Accessed: 2015-11-1.
- [2] Tor project. <https://www.torproject.org/>, 2015. Accessed: 2015-12-12.
- [3] R. Dingledine, N. Mathewson, and P. Syverson. Tor: The second-generation onion router. Technical report, DTIC Document, 2004.
- [4] D. Fisher. Link prefetching FAQ, 2003.
- [5] D. Fisher and G. Saksena. Link prefetching in mozilla: A server-driven approach. In *Web content caching and distribution*, pages 283–291. Springer, 2004.
- [6] M. Juarez, S. Afroz, G. Acar, C. Diaz, and R. Greenstadt. A critical evaluation of website fingerprinting attacks. In *Proceedings of the 2014 ACM SIGSAC Conference on Computer and Communications Security*, pages 263–274. ACM, 2014.
- [7] S. J. Murdoch and G. Danezis. Low-cost traffic analysis of tor. In *Security and Privacy, 2005 IEEE Symposium on*, pages 183–195. IEEE, 2005.
- [8] M. Nottingham. Rfc5988: Web linking. Technical report, 2010.
- [9] A. Panchenko, L. Niessen, A. Zinnen, and T. Engel. Website fingerprinting in onion routing based anonymization networks. In *Proceedings of the 10th annual ACM workshop on Privacy in the electronic society*, pages 103–114. ACM, 2011.