

# Link Prefetching: A Defense Against Website Fingerprinting on Tor \*

Vaibhav Sharma  
sharm361@umn.edu

Elaheh Ghassabani  
ghass013@umn.edu

Taejoon Byun  
taejoon@umn.edu

Se Eun Oh  
seoh@umn.edu

Department of Computer Science and Engineering  
University of Minnesota  
Minneapolis, MN 55454

## ABSTRACT

*This paper is written to get an A in the CSCI5271 course. PERIOD.*

*This content will be edited later.* We plan to explore the area of website fingerprinting in anonymization networks starting with the paper Website Fingerprinting in Onion Routing Based Anonymization Networks.

## Keywords

Website fingerprinting, anonymity, encrypted traffic, Tor

## 1. INTRODUCTION

Anonymizing networks are privacy technologies that provide a mechanism to anonymize internet communications so as to protect users from network eavesdroppers. Although such systems are able to hide the communication (including both routing information and content), an attacker is still able to obtain different information by analyzing the network traffic. Network analysis can provide very rich information about message length, timing, and frequency by which an attacker can easily identify the communicating parties, and therefore bypass an anonymizing system. This problem is known as Website Fingerprinting (WF) attack, where the adversary attempts to recognize the encrypted traffic patterns of specific web pages without using any other information [10, 12].

This report is organized as follows;

## 2. WEBSITE FINGERPRINTING ON TOR

This section explains the background about Tor, an anonymity network, and the website fingerprinting attack which can

neutralize the anonymity that Tor provides.

### 2.1 Tor

The Onion Router (Tor) is one of the most popular anonymizing networks, which distributes user's communications over several places on the Internet. The distribution helps user's final destination not to be linked to a single point. In other words, packets take a random path through several relays, instead of taking a direct route. These layers cover the user path such that it will not be possible to tell where the data comes from, or where it is going. Such a pathway in Tor is called a private network pathway. The Tor client application is responsible of creating this private pathway by incrementally building a circuit of encrypted connections. The circuit is constructed hop by hop, which means each individual relay is only aware of the previous and next hop. That is to say, individual relays never can identify the complete path [2, 4].

If an attacker can see both ends of the private pathway, Tor will fail. For example, suppose the attacker watches the Tor relay to which the user enters the network, and also watches the website he visits; In this case, the attacker can correlate volume and timing information on the two sides. To cope with this problem, Tor came up with the idea of "entry guards", where each Tor client randomly selects a few relays as entry points for her first hop. If those relays are not observed by an attacker, the user is secure. But, needless to say, there is always a probability of losing anonymity [2].

### 2.2 Website Fingerprinting

Tor does not provide protection against all anonymity problems. In addition WF attack is still able to bypass its privacy mechanism. In the case of Tor, WF attack would take place between the user and the Guard node, or at the Guard node itself. In general, attack models can be divided into two categories: "closed world" and "open world" scenarios. This section describes these attack models.

In the closed world model, the classifier can successfully recognize only web pages that it has already been trained on. Therefore, in this case, we deal with a small set of censored web pages. The open world scenario is when the ability of the adversary can recognize censored pages that might have not been seen before [1].

Figure 1 illustrates the attack model of website fingerprinting.

\*This report is submitted as a partial fulfillment of CSCI5271: Introduction to Security course.



Figure 1: Website Fingerprinting Attack Model

### 3. RELATED WORK

This section surveys attacks/ defenses on/against anonymizing systems, especially Tor.

#### 3.1 Fingerprinting Attacks

The idea of extracting information about the content of encrypted SSL packets dates back to 1996 [17]. Later in 2002, Hintz presented an attack on an encrypting web proxy calling it website fingerprinting. The attack successfully exploited a traffic analysis-based vulnerability to detect which website a particular user is surfing [9]. Since then, a lot of research has been conducted in this area trying either to propose a more realistic attack, or a more effective defense mechanism.

One interesting study has been done in [8], where authors launched WF attack on different anonymizing techniques, including OpenSSH, OpenVPN, Stunnel, Tor. They tried to attack 775 different websites. The result presents a low detection rate in the closed world attack model (only a 3% success rate for 775 pages). They used Multinomial Naive Bayes classifier focusing only on packet sizes. Later on, this attack has been improved with around 90% accuracy for 100 webpages in [18, 3]. In [3], authors could defeat ad hoc defense mechanisms describing a new defense scheme that provides provable security properties. Penchenko et al. [14] were also among the first to report website fingerprinting attacks with reasonable accuracy on Tor. They provided a sufficient understanding of the feature set and classification framework required for this attack. In fact, they carried out a comprehensive study of WF on Tor, which considers both closed world and open world attack models.

Generally speaking, attacks on anonymizing networks take a variety of approaches; some of the attacks tries to discover the identity of the anonymous user, others focus on uncovering the private pathway, and others attempt to identify the servers users interact with [3].

#### 3.2 Defense Mechanisms

Defense mechanisms are usually developed at the IP/TCP level by changing the pattern of traffic. For example, they split packets into multiple packets, insert fake packets into the traffic, or involve padding packets. A study performed in [7] shows that transmission at random intervals could be a defense against analysis-based attacks. Another defense technique proposed in [19], called traffic morphing, uses some tricks to change a traffic pattern in a way that it looks like another pattern. However, this method does not split or disordered packets. Therefore, attacks that work without packet size information can easily defeat this mechanism.

In [11], Luo et al. proposed a collection of tricks at the HTTP/TCP level as a defense mechanism against some analysis attacks. For example, they changed window sizes and the order of packets in the TCP stream. Besides, at the

HTTP level, they tried to insert some extra data into HTTP GET headers, generate some irrelevant HTTP request, and re-order requests. The Tor community also introduced "randomized pipelining" [15] as a defense mechanism, where the browser loads the web content in a random order. Despite such techniques, the proposed attack in [3] managed to successfully recognize target web pages with the accuracy of 87%. The experiments show that techniques like traffic morphing and randomized pipelining can impose high costs, but are unable to stop the WF attack. This attack is able to ignore packet sizes, while most of the attacks on Tor work based on packet sizes.

#### 3.3 Criticism

In the literature, some work always debate over the practical feasibility of the WF attack. This section tries to summarize issues questioning the practicality of WF.

There are many different factors that play a role in the success of WF. The main criticism is that academic papers usually oversimplify WF attack models by making some unrealistic assumptions over users' browsing habits, training and testing traces, even the version of Tor used for testing/training. In [10], it is discussed that the studies performed in [3, 8, 14, 18, 16] simplify the problem and overestimate the adversary's capabilities.

Another criticism is that all works launch attacks on individual pages, instead of overall websites. So, although the attack is called website fingerprinting, it is actually about webpage fingerprinting. In addition, there are some main factors, including the hypothesis state space and the size of the instance space, that affect the accuracy of a classifier. Even, the number of training samples provided to the classifier and false positive rates matters. Since reliable feature information is constant, with the increase in the number of classification categories, the classifier eventually runs out of descriptive feature information, which causes either true positive accuracy goes down or the false positive rate goes up. [1] discusses that the effects of such factors are quite observable in the papers with a sufficient world size.

Although more attention should be paid to the scenarios by which attacks are evaluated, we should not dismiss WF as a threat. However, [1] argues that, due to theoretical and practical issues, realistic WF attacks are hard to launch on Tor. Therefore, it claims that even simple defenses could protect Tor users against WF. It seems that the Tor community believes that defense mechanisms do not need to be very complicated to be effective.

To the best of our knowledge, none of the defense techniques has investigated the effect of link prefetching on WF. In the next section, we will describe the concept of link prefetching. Since most of the defenses are applied at the Tor network, they are required to be acceptable by the Tor community. However, we aim to propose a defense technique which can also be used by the website owners.

### 4. LINK PREFETCHING AS A DEFENSE

This section explains the concept of link prefetching and discusses a possible usage of it as a defense mechanism against fingerprinting attacks on Tor.

#### 4.1 Link Prefetching

Link prefetching is a HTML syntax that gives the web browser hints about which page the user is most likely to

visit in a near future [5, 6]. The pages and resource to pre-fetch are specified in the web page so that the web browser can silently load them (or pre-fetch them) after an idle time. Since the pre-fetch happens only after the page is fully loaded, it does not sacrifice the loading time of the requested web page. Moreover, it can save the loading time for the pre-fetched pages and thus improve the user experience by caching the *future* contents. It was first suggested by Mozilla Foundation in 2003 and supported by most modern browsers nowadays.

The resources to pre-fetch can be simply specified in *HTML* using a `link` tag [13]. For example, a `link` tag `<link rel="prefetch" href="/page2.html">` tells the browser to pre-fetch a *html* file named *page2.html*. Resources other than a *HTML* web page can also be pre-fetched similarly using the same syntax. There are also some variations for different types of prefetching – DNS prefetching, which is specified as `<link rel="dns-prefetch" .>`, is supported by *Mozilla Firefox* and *Google Chrome*, and `<link rel="prerender" .>` also does the same job as `prefetch` in *Google Chrome* and *Microsoft Internet Explorer*.

Figure 2 illustrates how prefetching actually works in a browser (*Google Chrome*). This page is set up arbitrarily by the authors to demonstrate link prefetching, and contains a link pre-fetch tag that specifies to load a big image (the image on the left side labeled as *prefetch*). When the prefetching is off, this image shall be requested only when a user puts his mouse cursor on it (mouse over). However, it can be seen on the network timeline (right bottom) that the image is pre-fetched right after loading the page. This is indicated by a long blue bar on the second row for the file named *Very-high...*, and it is long because the size of the file is relatively big (3.5 MB) that it took a relatively long time to download. Please also note that the time it took for loading the image when the user actually requested (by putting his mouse on it) was very short, because the image had already been pre-fetched that the browser merely loaded it from the cache (as shown in the *size* column on the fourth row).

## 4.2 The Effect of Prefetching on Fingerprinting

The idea of studying the effect of prefetching on Tor network is not new [?].

Figure 3 shows an aspect of fingerprint we gathered for the first page of [www.wired.com](http://www.wired.com). It depicts the number of packets on the *y-axis* for each time interval (500 ms), where the *x-axis* represents the time in seconds. The solid line represents the number of packets throughout time when the pre-fetch feature was enabled, and the dashed line represents when the pre-fetch was disabled in the browser settings.

## 5. EXPERIMENTAL EVALUATION

Experimental result comes here.

### 5.1 Research Question

1. **RQ1:** Does prefetching itself provide an extra degree of defense?
2. **RQ2:** Can prefetching be used as a browser-side defense mechanism?

**Table 1: Experiment design to answer RQ1**

victim \ attacker	prefetch on	prefetch off
prefetch on	(1)	(2)
prefetch off	(3)	(4)

1. We speculate that prefetching itself might provide extra defense because of the extra packets. It can also be the case however, prefetching websites are more vulnerable to fingerprinting because of the extra prefetch packets that shows distinct prefix.
2. This case is unlikely since prefetching is on by default. We assume that victims will more likely be using Tor under the default setting.
3. This case is what we are most curious about, whether a victim can confuse an attacker by simply turning prefetching setting off of his browser.
4. This case simulates a situation where a victim is loading any other websites that does not prefetch any resource. This can be used as a comparison case.

### 5.2 Effects of Pre-Fetching on Fingerprinting

In this section, we will write about our experiments. We are planning to conduct two sets of experiments. If we consider the network traffic, the number of packages go upstream depends on the number of pre-fetching requests, and the number of downstream packages coming depends on the size of resources that should be pre-fetched. Therefore, it is obvious that pre-fetching would affect the fingerprint of the traffic of a particular website. *should be completed*.

### 5.3 Investigate Pre-Fetching Effects on top 60 Popular Websites

We are running experiments to see how pre-fetching affects the websites' fingerprints. After doing some search on top popular websites, we put together a small crawler by which we learnt that only around 60 of all 6000 websites are use pre-fetching mechanism. We are capturing traffic of these websites in two different modes: 1) with enabled prefetching, and 2) with disabled pre-fetching. We are working on feature extraction, and about to decide which classifiers to use for the learning phase. Ultimately, we plan to conduct two sets of experiments. One sort of experiment is to compare two series of the captured packets and find the accuracy number with the help of a classifier, by which our goal is to provide an evidence to see if pre-fetching really affects fingerprints of websites. So, if the result will be positive, we will perform another set of experiment, which kind of simulates a sub set of those 60 websites. Then, we will see how (altering) the size of pre-fetching affects fingerprinting attacks/ defense mechanisms.

### 5.4 Effect of Pre-Fetching Packets Size on Fingerprinting Attacks

Here, we will explain our second experiment. We will simulate a sub set of webpages we investigated in the previous experiment. Then, we will equip them with a mechanism so that they can affect the downstream traffic and finally their fingerprint. Then, we will analyze the result to see how this




Figure 2: Network timeline showing pre-fetch




Figure 3: hi

idea contributes to the effectiveness of attacks and defense techniques.

## 6. DISCUSSION

We'll find what to discuss later on after finishing the experiment section.

## 7. CONCLUSIONS

What we have done in this work is so cool and awesome. What you may criticize will all be put here as "future work".

This section will conclude the result of our experiments. Finally we will provide some evidence to show how pre-fetching affect fingerprinting attacks. Based on our result, we are planing to suggest some defense mechanisms.

## 8. ACKNOWLEDGMENTS

This is a research project for CSCI5271, University of Minnesota.

## 9. REFERENCES

- [1] A critique of website traffic fingerprinting attacks. <https://blog.torproject.org/blog/critique-website-traffic-fingerprinting-attacks>, 2015. Accessed: 2015-11-1.
- [2] Tor project. <https://www.torproject.org/>, 2015. Accessed: 2015-12-12.
- [3] X. Cai, X. C. Zhang, B. Joshi, and R. Johnson. Touching from a distance: Website fingerprinting attacks and defenses. In *Proceedings of the 2012 ACM conference on Computer and communications security*, pages 605–616. ACM, 2012.
- [4] R. Dingledine, N. Mathewson, and P. Syverson. Tor: The second-generation onion router. Technical report, DTIC Document, 2004.
- [5] D. Fisher. Link prefetching FAQ. [https://developer.mozilla.org/en-US/docs/Web/HTTP/Link\\_prefetching\\_FAQ](https://developer.mozilla.org/en-US/docs/Web/HTTP/Link_prefetching_FAQ), 2003.

- [6] D. Fisher and G. Saksena. Link prefetching in mozilla: A server-driven approach. In *Web content caching and distribution*, pages 283–291. Springer, 2004.
- [7] X. Fu, B. Graham, R. Bettati, and W. Zhao. On countermeasures to traffic analysis attacks. In *Information Assurance Workshop, 2003. IEEE Systems, Man and Cybernetics Society*, pages 188–195. IEEE, 2003.
- [8] D. Herrmann, R. Wendolsky, and H. Federrath. Website fingerprinting: attacking popular privacy enhancing technologies with the multinomial naïve-bayes classifier. In *Proceedings of the 2009 ACM workshop on Cloud computing security*, pages 31–42. ACM, 2009.
- [9] A. Hintz. Fingerprinting websites using traffic analysis. In *Privacy Enhancing Technologies*, pages 171–178. Springer, 2003.
- [10] M. Juarez, S. Afroz, G. Acar, C. Diaz, and R. Greenstadt. A critical evaluation of website fingerprinting attacks. In *Proceedings of the 2014 ACM SIGSAC Conference on Computer and Communications Security*, pages 263–274. ACM, 2014.
- [11] X. Luo, P. Zhou, E. W. Chan, W. Lee, R. K. Chang, and R. Perdisci. Https: Sealing information leaks with browser-side obfuscation of encrypted flows. In *NDSS*, 2011.
- [12] S. J. Murdoch and G. Danezis. Low-cost traffic analysis of tor. In *Security and Privacy, 2005 IEEE Symposium on*, pages 183–195. IEEE, 2005.
- [13] M. Nottingham. Rfc5988: Web linking. Technical report, 2010.
- [14] A. Panchenko, L. Niessen, A. Zinnen, and T. Engel. Website fingerprinting in onion routing based anonymization networks. In *Proceedings of the 10th annual ACM workshop on Privacy in the electronic society*, pages 103–114. ACM, 2011.
- [15] M. Perry. Experimental defense for website traffic fingerprinting. <https://blog.torproject.org/blog/experimental-defense-website-traffic-fingerprinting>. Accessed: 2015-11-23.
- [16] Y. Shi and K. Matsuura. Fingerprinting attack on the tor anonymity system. In *Information and Communications Security*, pages 425–438. Springer, 2009.
- [17] D. Wagner, B. Schneier, et al. Analysis of the ssl 3.0 protocol. In *The Second USENIX Workshop on Electronic Commerce Proceedings*, pages 29–40, 1996.
- [18] T. Wang and I. Goldberg. Improved website fingerprinting on tor. In *Proceedings of the 12th ACM workshop on Workshop on privacy in the electronic society*, pages 201–212. ACM, 2013.
- [19] C. V. Wright, S. E. Coull, and F. Monrose. Traffic morphing: An efficient defense against statistical traffic analysis. In *NDSS*, 2009.